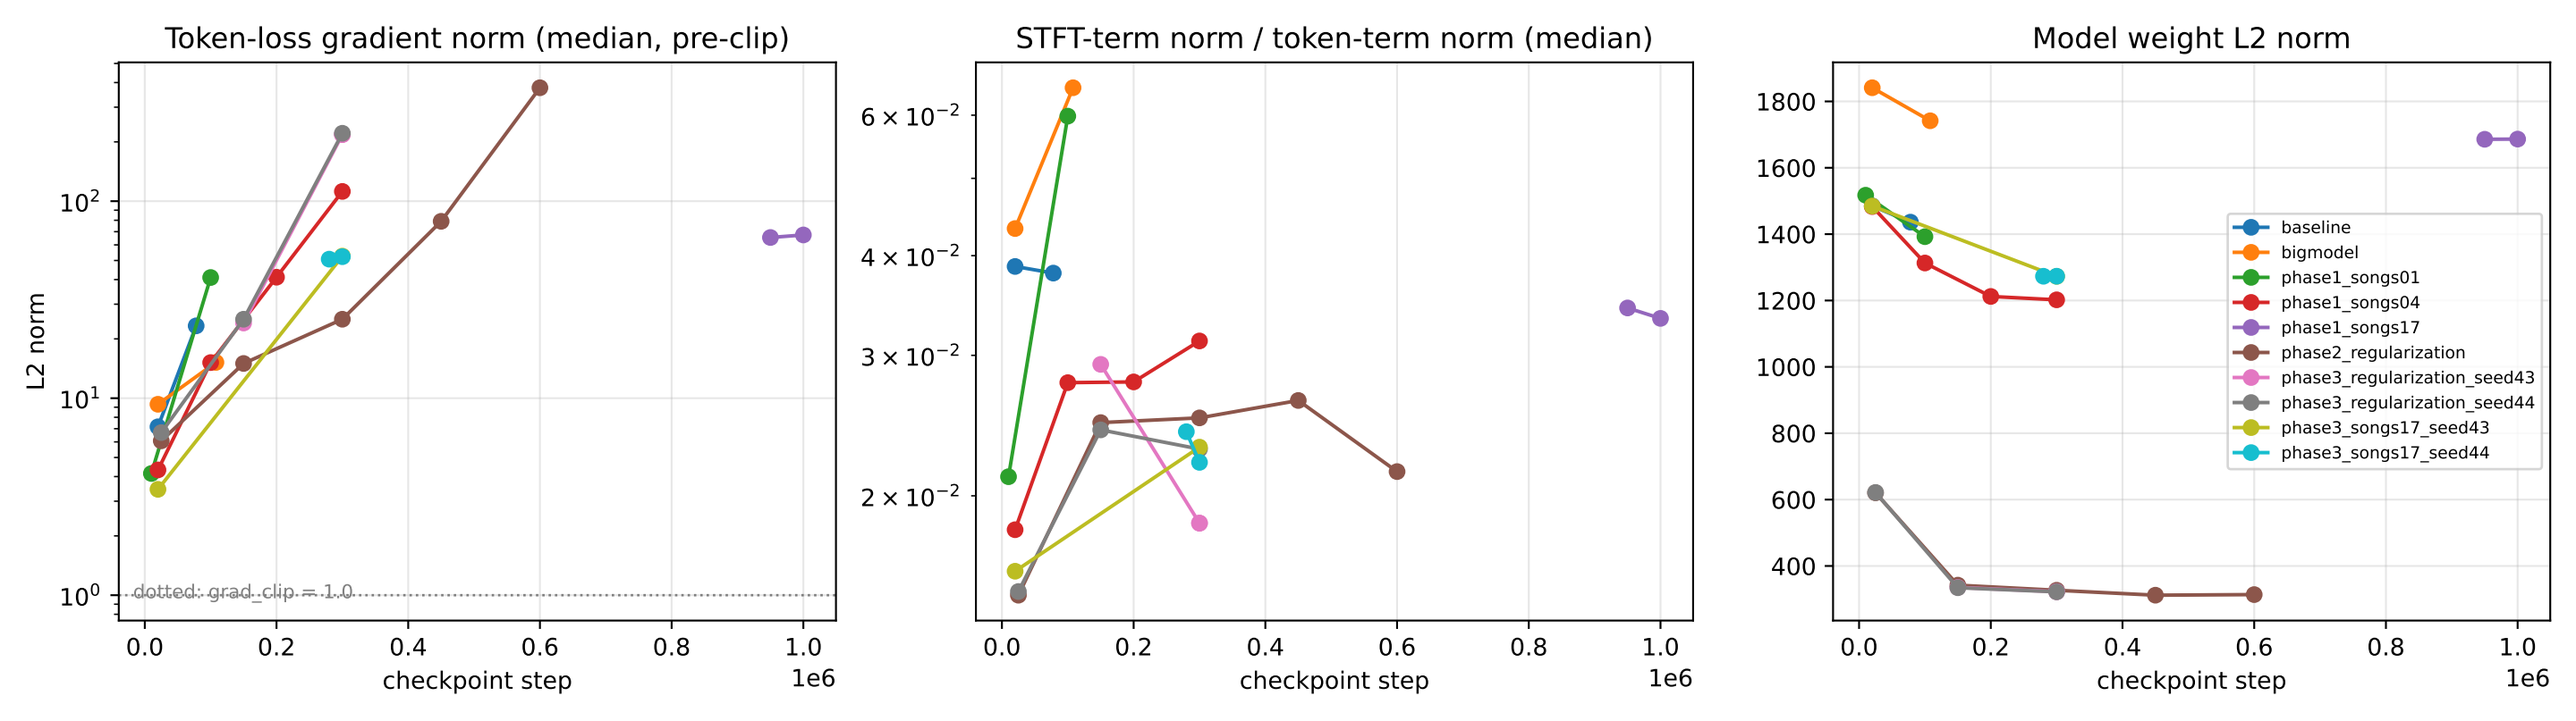
Token-loss gradient norm (median, pre-clip)
L2 norm
checkpoint step
1e6
dotted: grad_clip = 1.0
STFT-term norm / token-term norm (median)
checkpoint step
1e6
Model weight L2 norm
checkpoint step
1e6
baseline
bigmodel
phase1_songs01
phase1_songs04
phase1_songs17
phase2_regularization
phase3_regularization_seed43
phase3_regularization_seed44
phase3_songs17_seed43
phase3_songs17_seed44