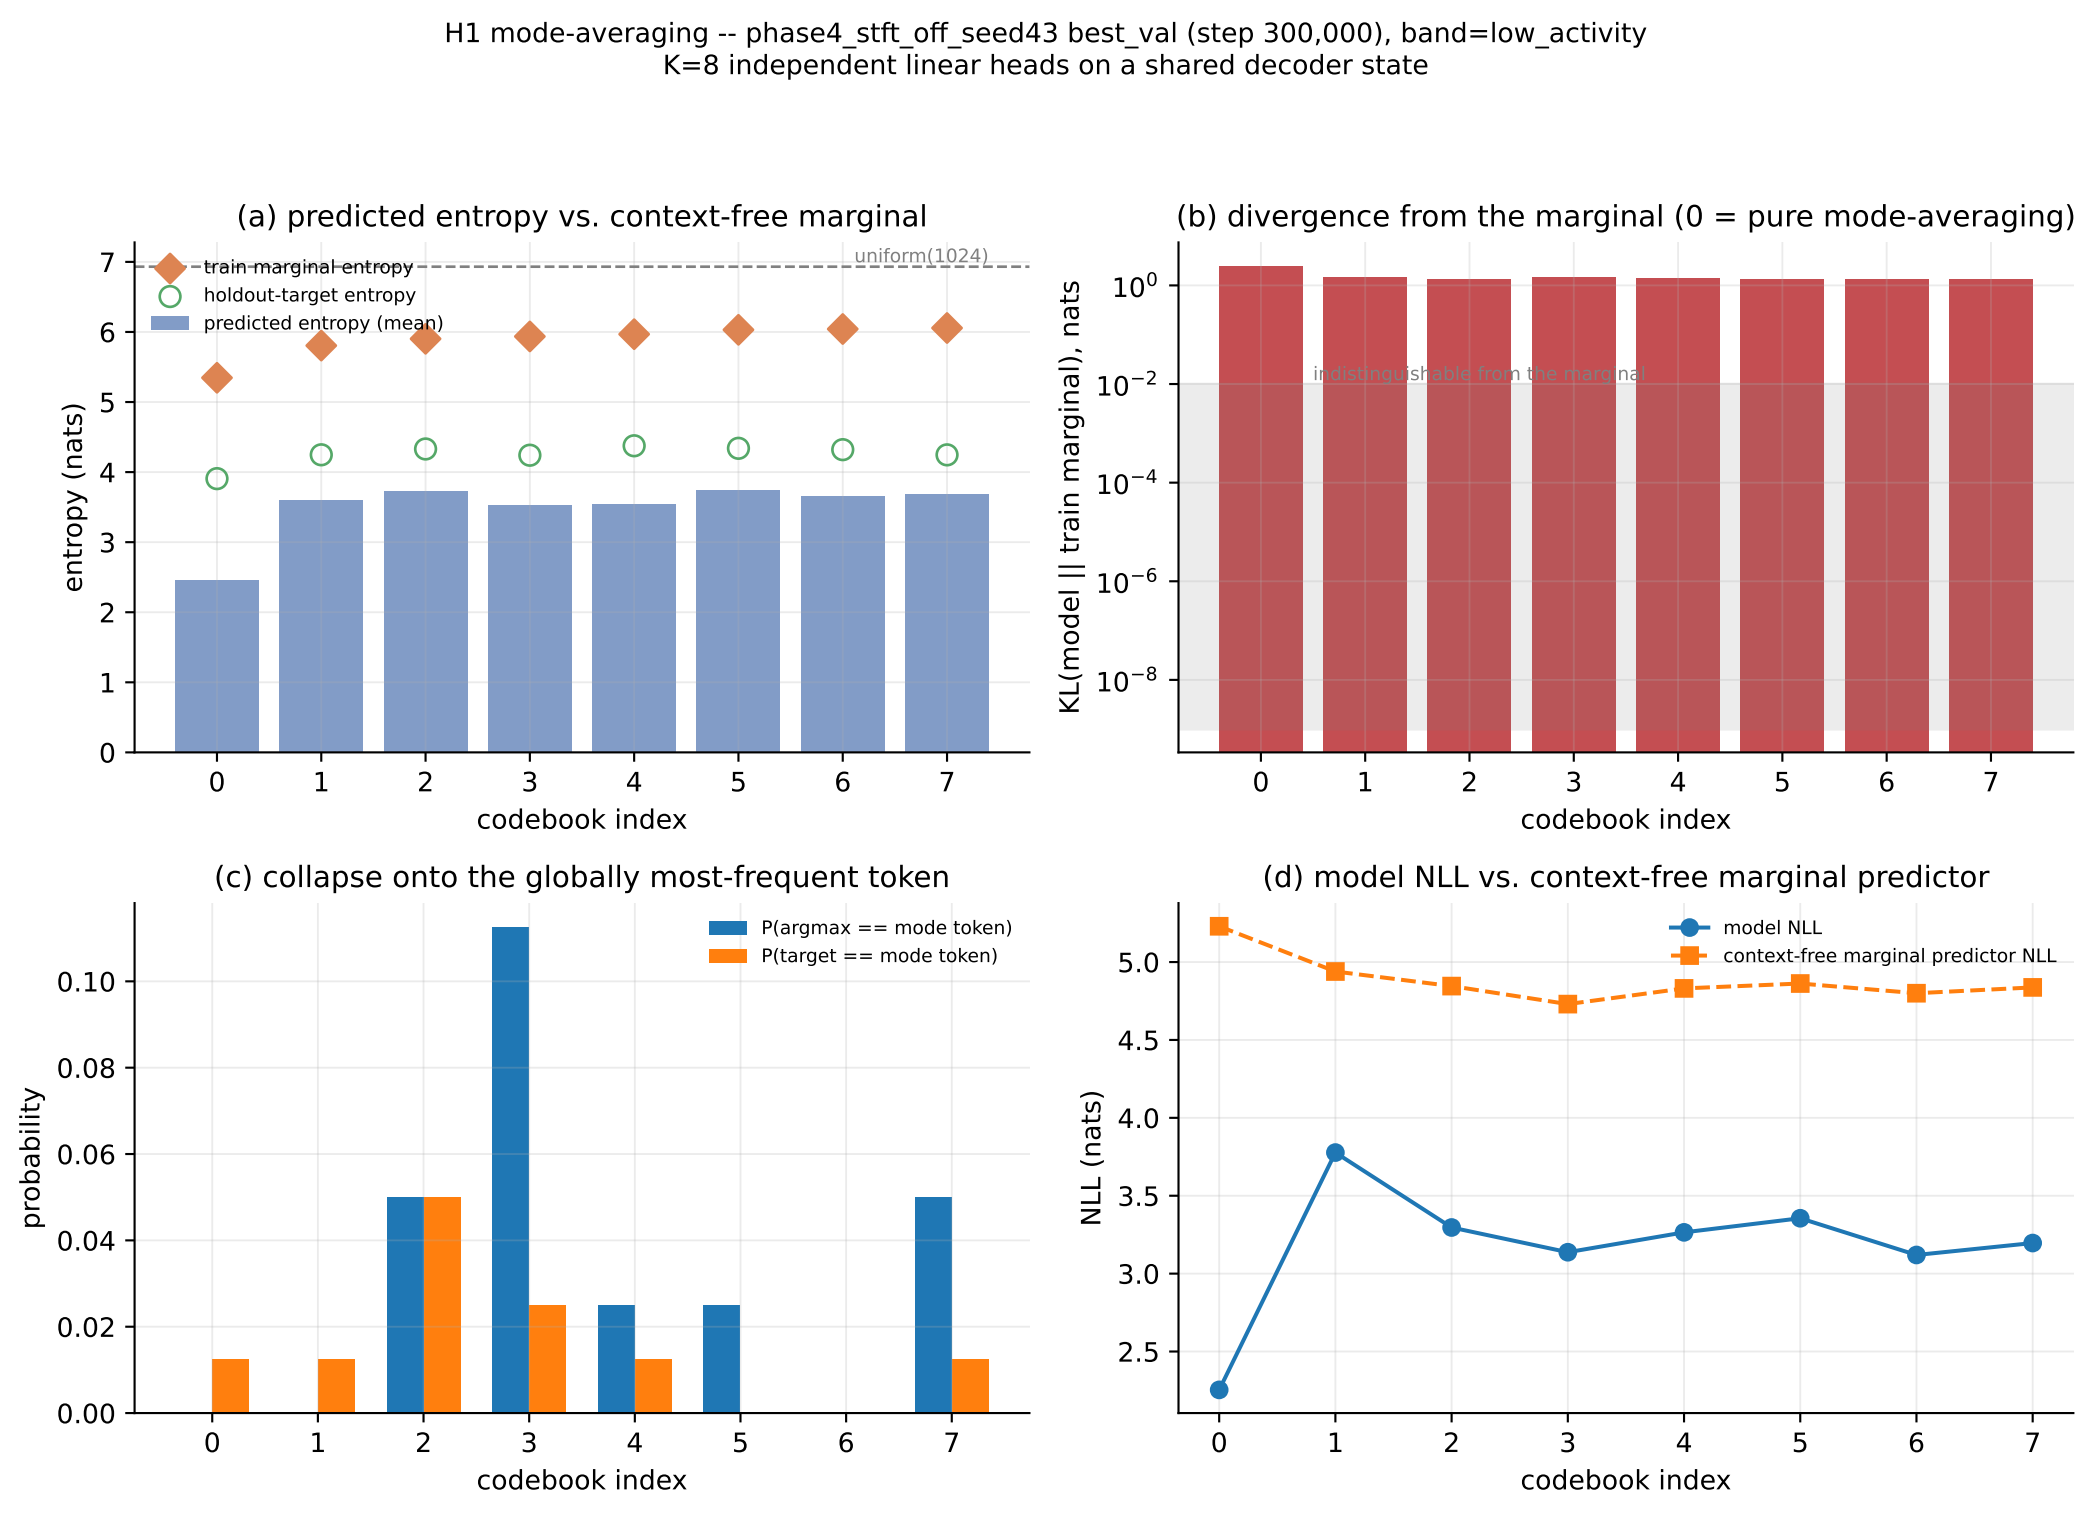
H1 mode-averaging -- phase4_stft_off_seed43 best_val (step 300,000), band=low_activity
K=8 independent linear heads on a shared decoder state
(a) predicted entropy vs. context-free marginal
train marginal entropy
holdout-target entropy
predicted entropy (mean)
uniform(1024)
entropy (nats)
codebook index
(b) divergence from the marginal (0 = pure mode-averaging)
KL(model || train marginal), nats
indistinguishable from the marginal
codebook index
(c) collapse onto the globally most-frequent token
P(argmax == mode token)
P(target == mode token)
probability
codebook index
(d) model NLL vs. context-free marginal predictor
model NLL
context-free marginal predictor NLL
NLL (nats)
codebook index