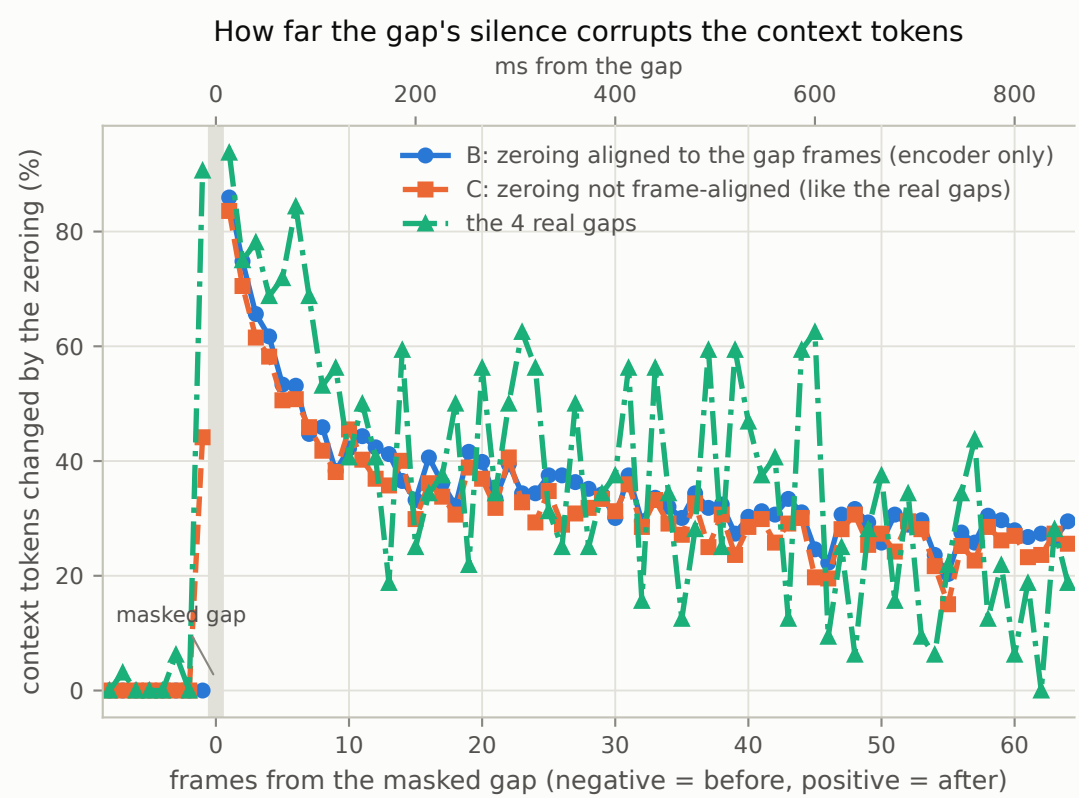
How far the gap's silence corrupts the context tokens
ms from the gap
0
200
400
600
800
B: zeroing aligned to the gap frames (encoder only)
C: zeroing not frame-aligned (like the real gaps)
the 4 real gaps
context tokens changed by the zeroing (%)
0
20
40
60
80
masked gap
0
10
20
30
40
50
60
frames from the masked gap (negative = before, positive = after)

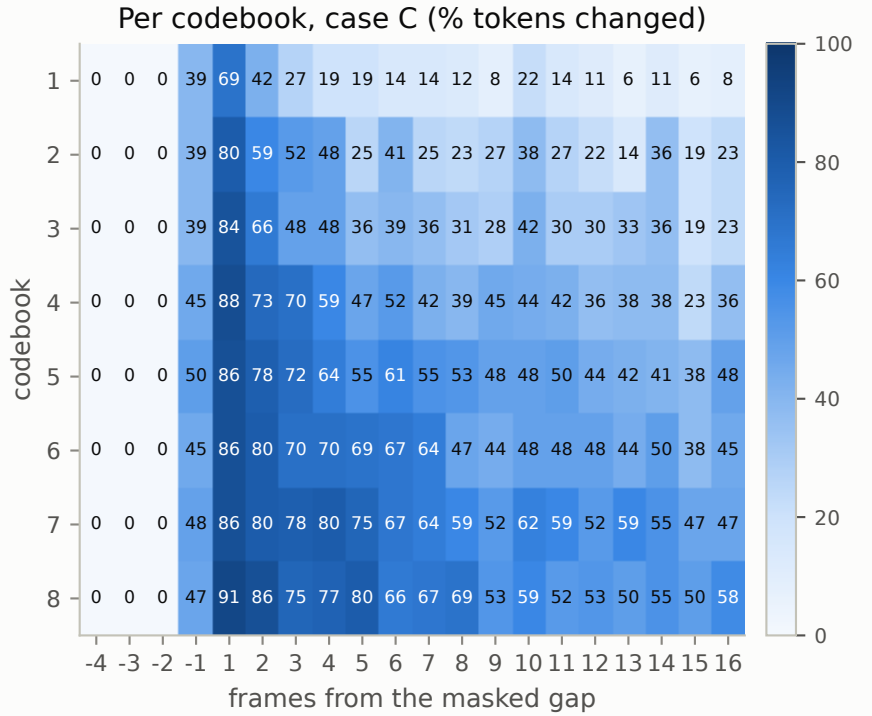
Per codebook, case C (% tokens changed)
codebook
1 0 0 0 39 69 42 27 19 19 14 14 12 8 22 14 11 6 11 6 8
2 0 0 0 39 80 59 52 48 25 41 25 23 27 38 27 22 14 36 19 23
3 0 0 0 39 84 66 48 48 36 39 36 31 28 42 30 30 33 36 19 23
4 0 0 0 45 88 73 70 59 47 52 42 39 45 44 42 36 38 38 23 36
5 0 0 0 50 86 78 72 64 55 61 55 53 48 48 50 44 42 41 38 48
6 0 0 0 45 86 80 70 70 69 67 64 47 44 48 48 48 44 50 38 45
7 0 0 0 48 86 80 78 80 75 67 64 59 52 62 59 52 59 55 47 47
8 0 0 0 47 91 86 75 77 80 66 67 69 53 59 52 53 50 55 50 58
-4 -3 -2 -1 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
frames from the masked gap
100
80
60
40
20
0