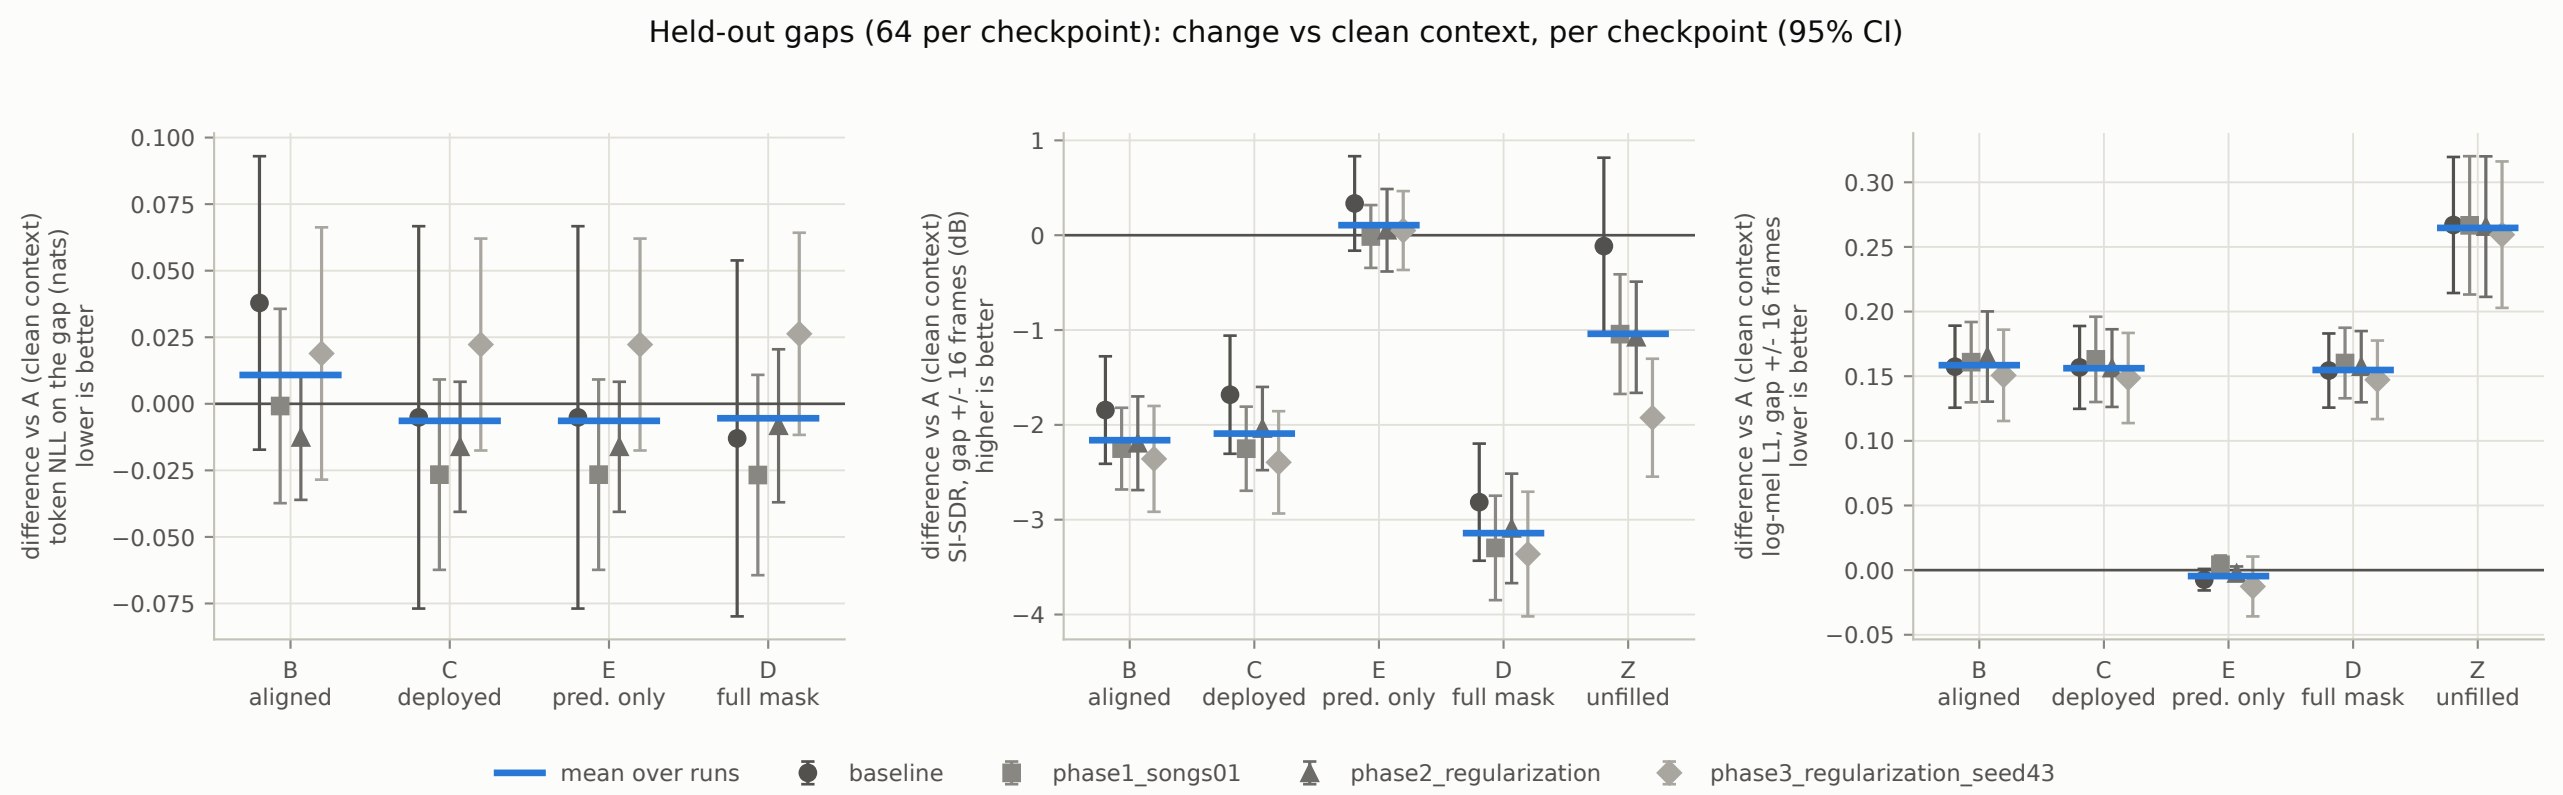
Held-out gaps (64 per checkpoint): change vs clean context, per checkpoint (95% CI)
difference vs A (clean context)
token NLL on the gap (nats)
lower is better
0.100
0.075
0.050
0.025
0.000
−0.025
−0.050
−0.075
B
aligned
C
deployed
E
pred. only
D
full mask
difference vs A (clean context)
SI-SDR, gap +/- 16 frames (dB)
higher is better
1
0
−1
−2
−3
−4
B
aligned
C
deployed
E
pred. only
D
full mask
Z
unfilled
difference vs A (clean context)
log-mel L1, gap +/- 16 frames
lower is better
0.30
0.25
0.20
0.15
0.10
0.05
0.00
−0.05
B
aligned
C
deployed
E
pred. only
D
full mask
Z
unfilled
mean over runs
baseline
phase1_songs01
phase2_regularization
phase3_regularization_seed43
A clean context (training-like) B frame-aligned zeroing C deployed today (unaligned zeroing, rounded mask)
E C's predictions written into clean context D mask covers every silent frame Z gap not inpainted