Real gaps (4 per checkpoint): change vs clean context, per checkpoint (95% CI)

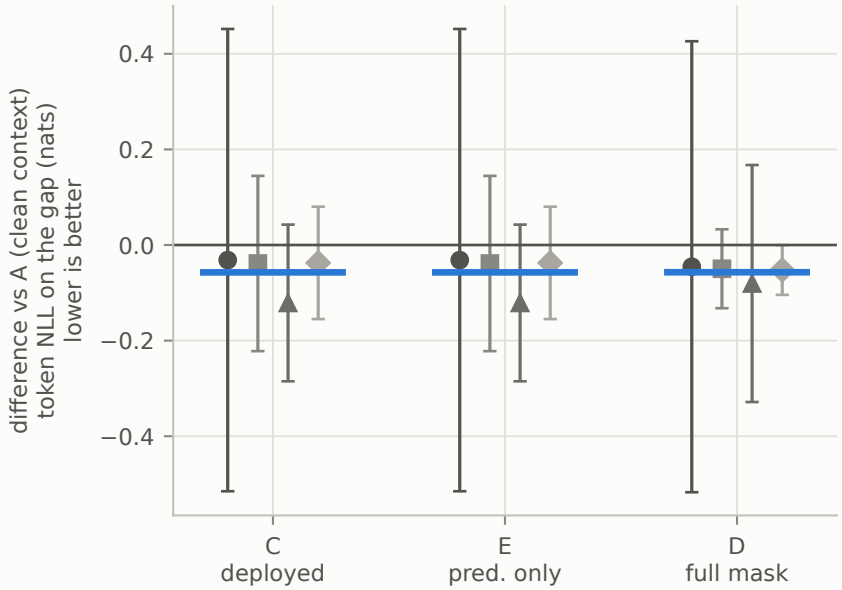

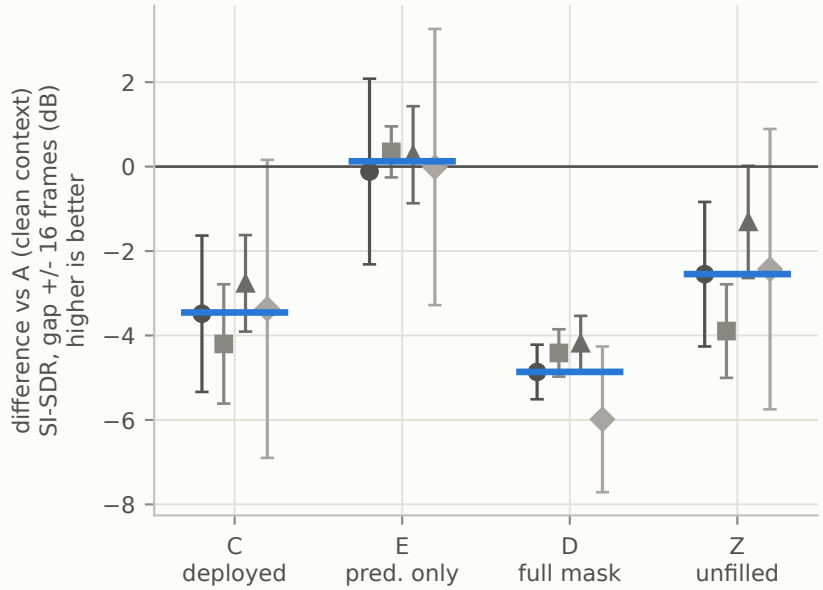

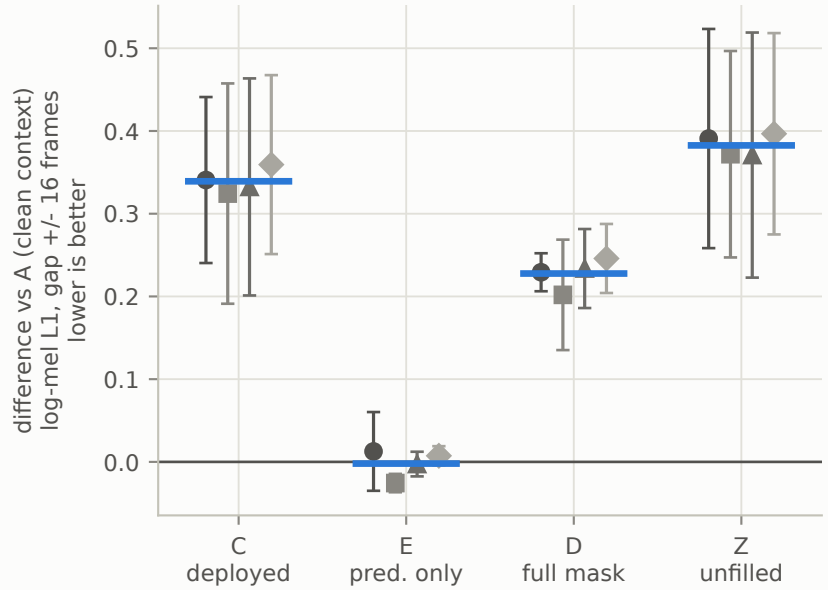

mean over runs baseline phase1_songs01 phase2_regularization phase3_regularization_seed43

A clean context (training-like) B frame-aligned zeroing C deployed today (unaligned zeroing, rounded mask)
E C's predictions written into clean context D mask covers every silent frame Z gap not inpainted