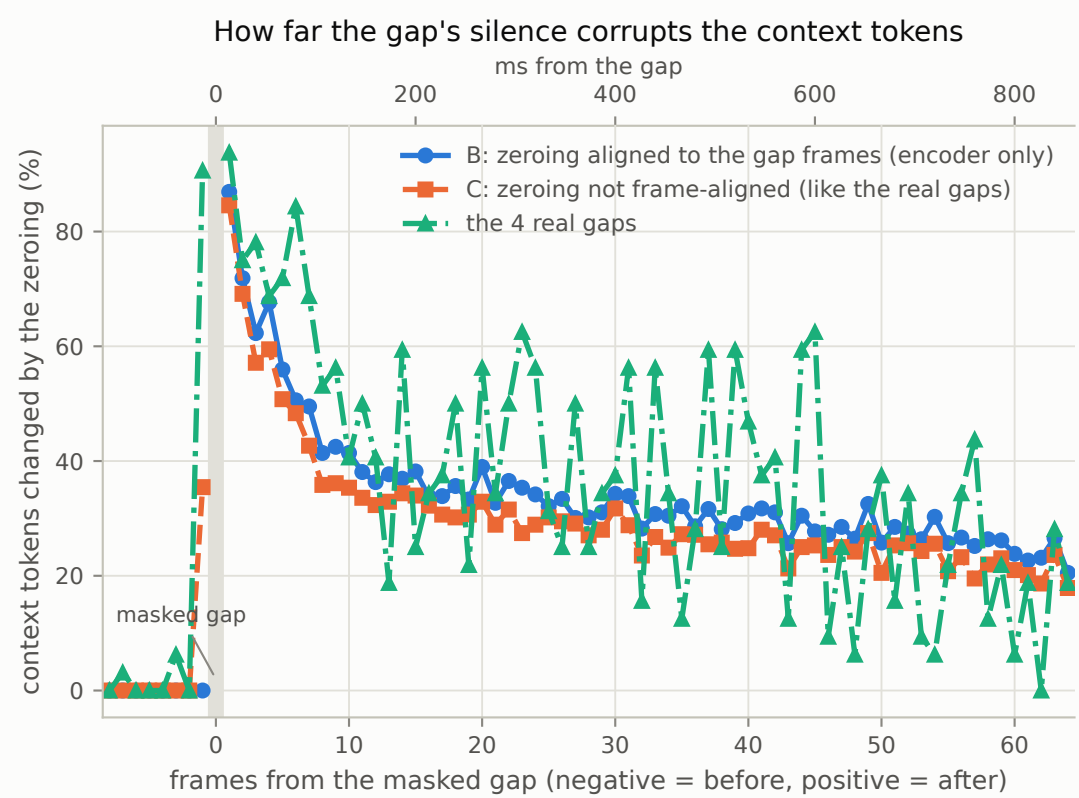
How far the gap's silence corrupts the context tokens
ms from the gap
0
200
400
600
800
B: zeroing aligned to the gap frames (encoder only)
C: zeroing not frame-aligned (like the real gaps)
the 4 real gaps
masked gap
context tokens changed by the zeroing (%)
0
20
40
60
80
0
10
20
30
40
50
60
frames from the masked gap (negative = before, positive = after)

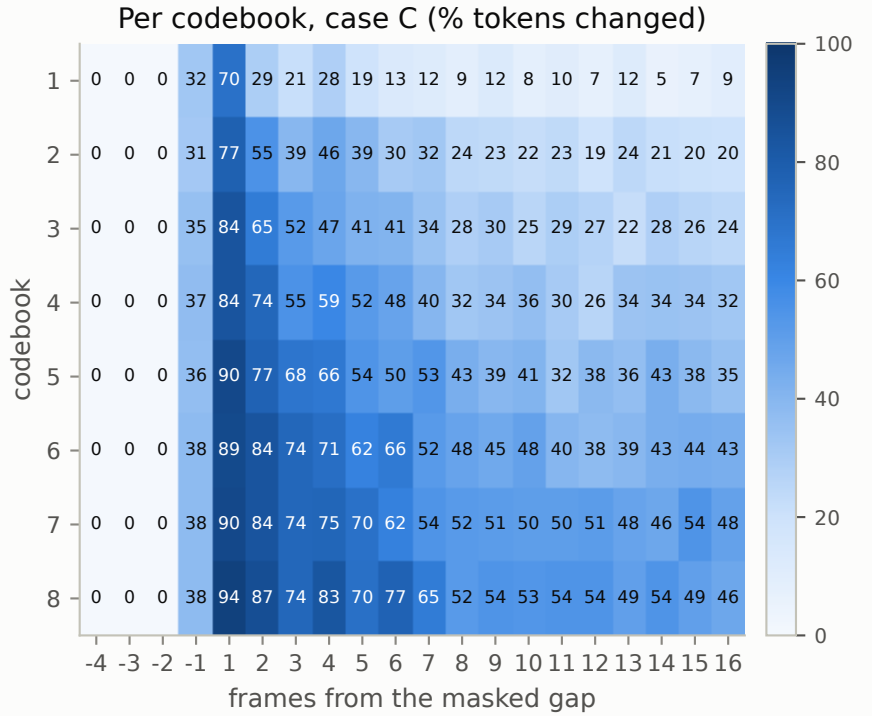
Per codebook, case C (% tokens changed)
codebook
1 0 0 0 32 70 29 21 28 19 13 12 9 12 8 10 7 12 5 7 9
2 0 0 0 31 77 55 39 46 39 30 32 24 23 22 23 19 24 21 20 20
3 0 0 0 35 84 65 52 47 41 41 34 28 30 25 29 27 22 28 26 24
4 0 0 0 37 84 74 55 59 52 48 40 32 34 36 30 26 34 34 34 32
5 0 0 0 36 90 77 68 66 54 50 53 43 39 41 32 38 36 43 38 35
6 0 0 0 38 89 84 74 71 62 66 52 48 45 48 40 38 39 43 44 43
7 0 0 0 38 90 84 74 75 70 62 54 52 51 50 50 51 48 46 54 48
8 0 0 0 38 94 87 74 83 70 77 65 52 54 53 54 54 49 54 49 46
-4 -3 -2 -1 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
frames from the masked gap
100
80
60
40
20
0