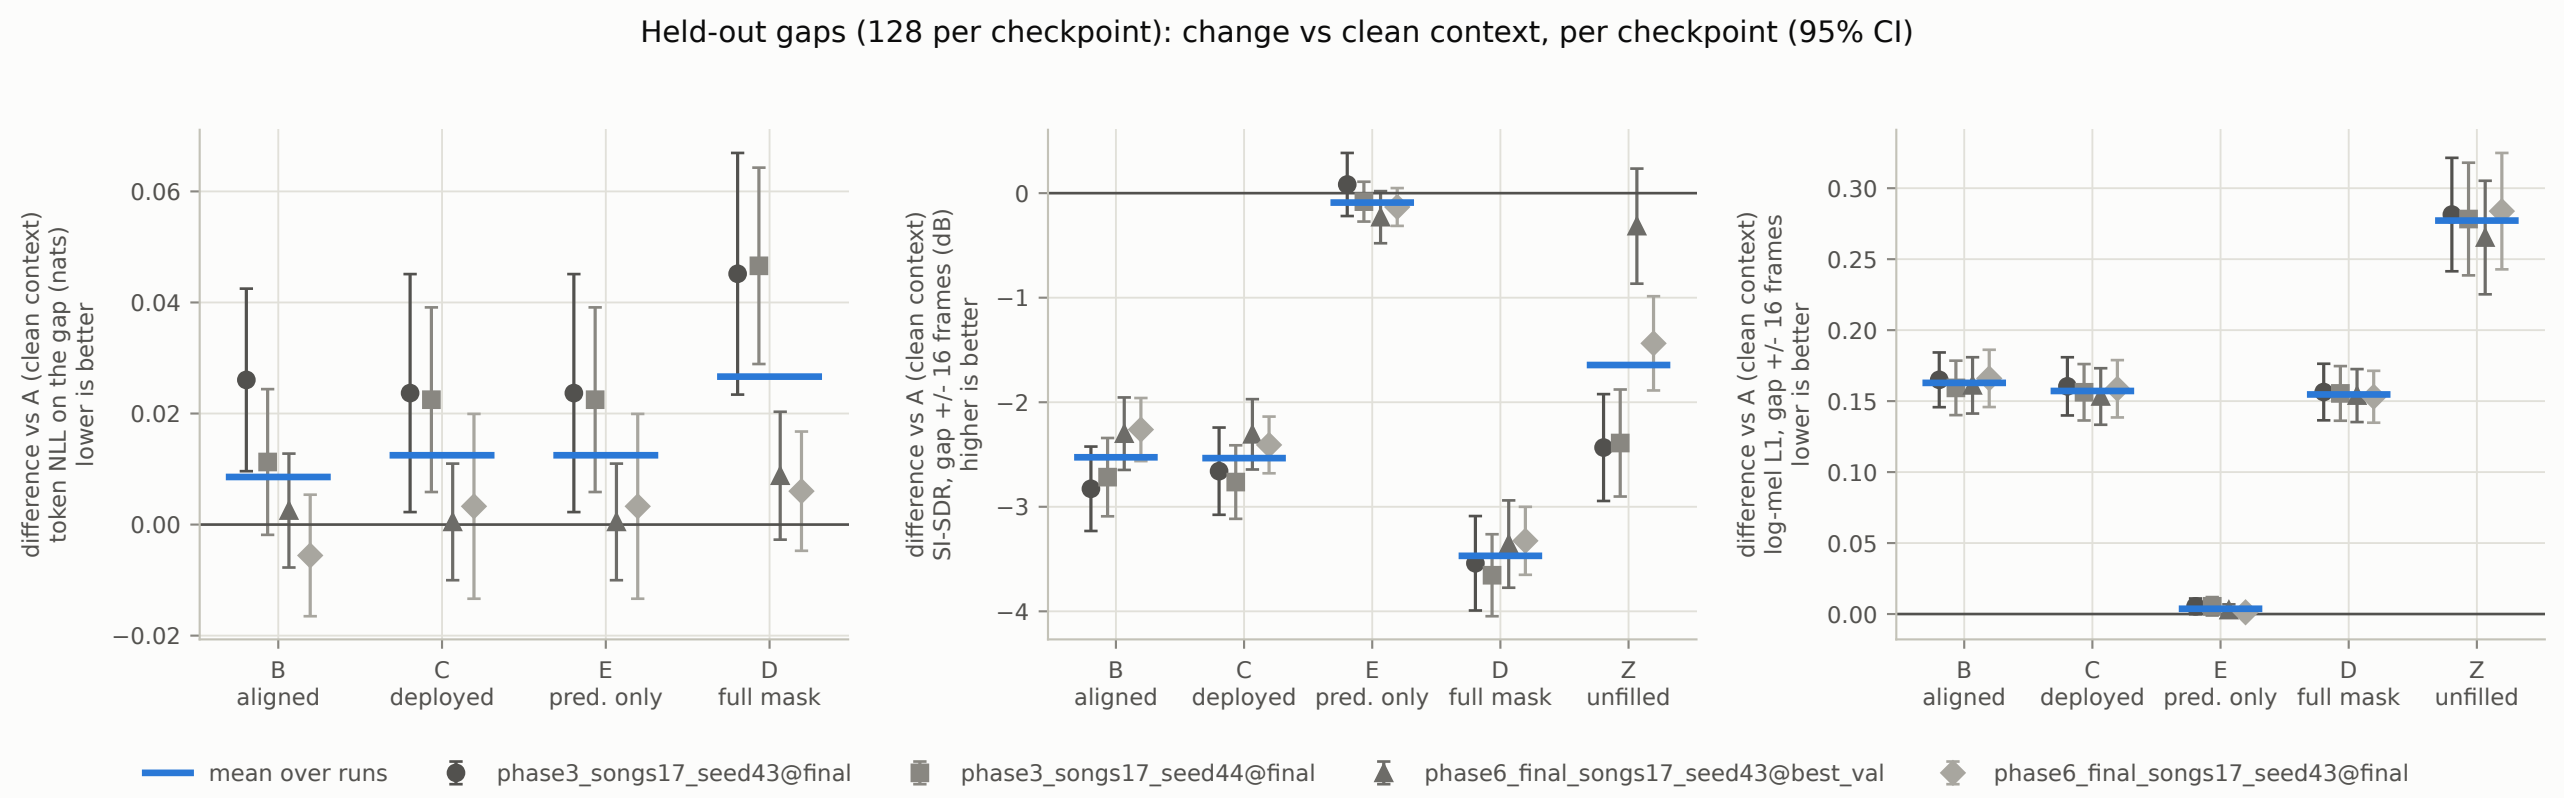
Held-out gaps (128 per checkpoint): change vs clean context, per checkpoint (95% CI)
difference vs A (clean context)
token NLL on the gap (nats)
lower is better
0.06
0.04
0.02
0.00
−0.02
B
aligned
C
deployed
E
pred. only
D
full mask
difference vs A (clean context)
SI-SDR, gap +/- 16 frames (dB)
higher is better
0
−1
−2
−3
−4
B
aligned
C
deployed
E
pred. only
D
full mask
Z
unfilled
difference vs A (clean context)
log-mel L1, gap +/- 16 frames
lower is better
0.30
0.25
0.20
0.15
0.10
0.05
0.00
B
aligned
C
deployed
E
pred. only
D
full mask
Z
unfilled
mean over runs
phase3_songs17_seed43@final
phase3_songs17_seed44@final
phase6_final_songs17_seed43@best_val
phase6_final_songs17_seed43@final
A clean context (training-like) B frame-aligned zeroing C deployed today (unaligned zeroing, rounded mask)
E C's predictions written into clean context D mask covers every silent frame Z gap not inpainted