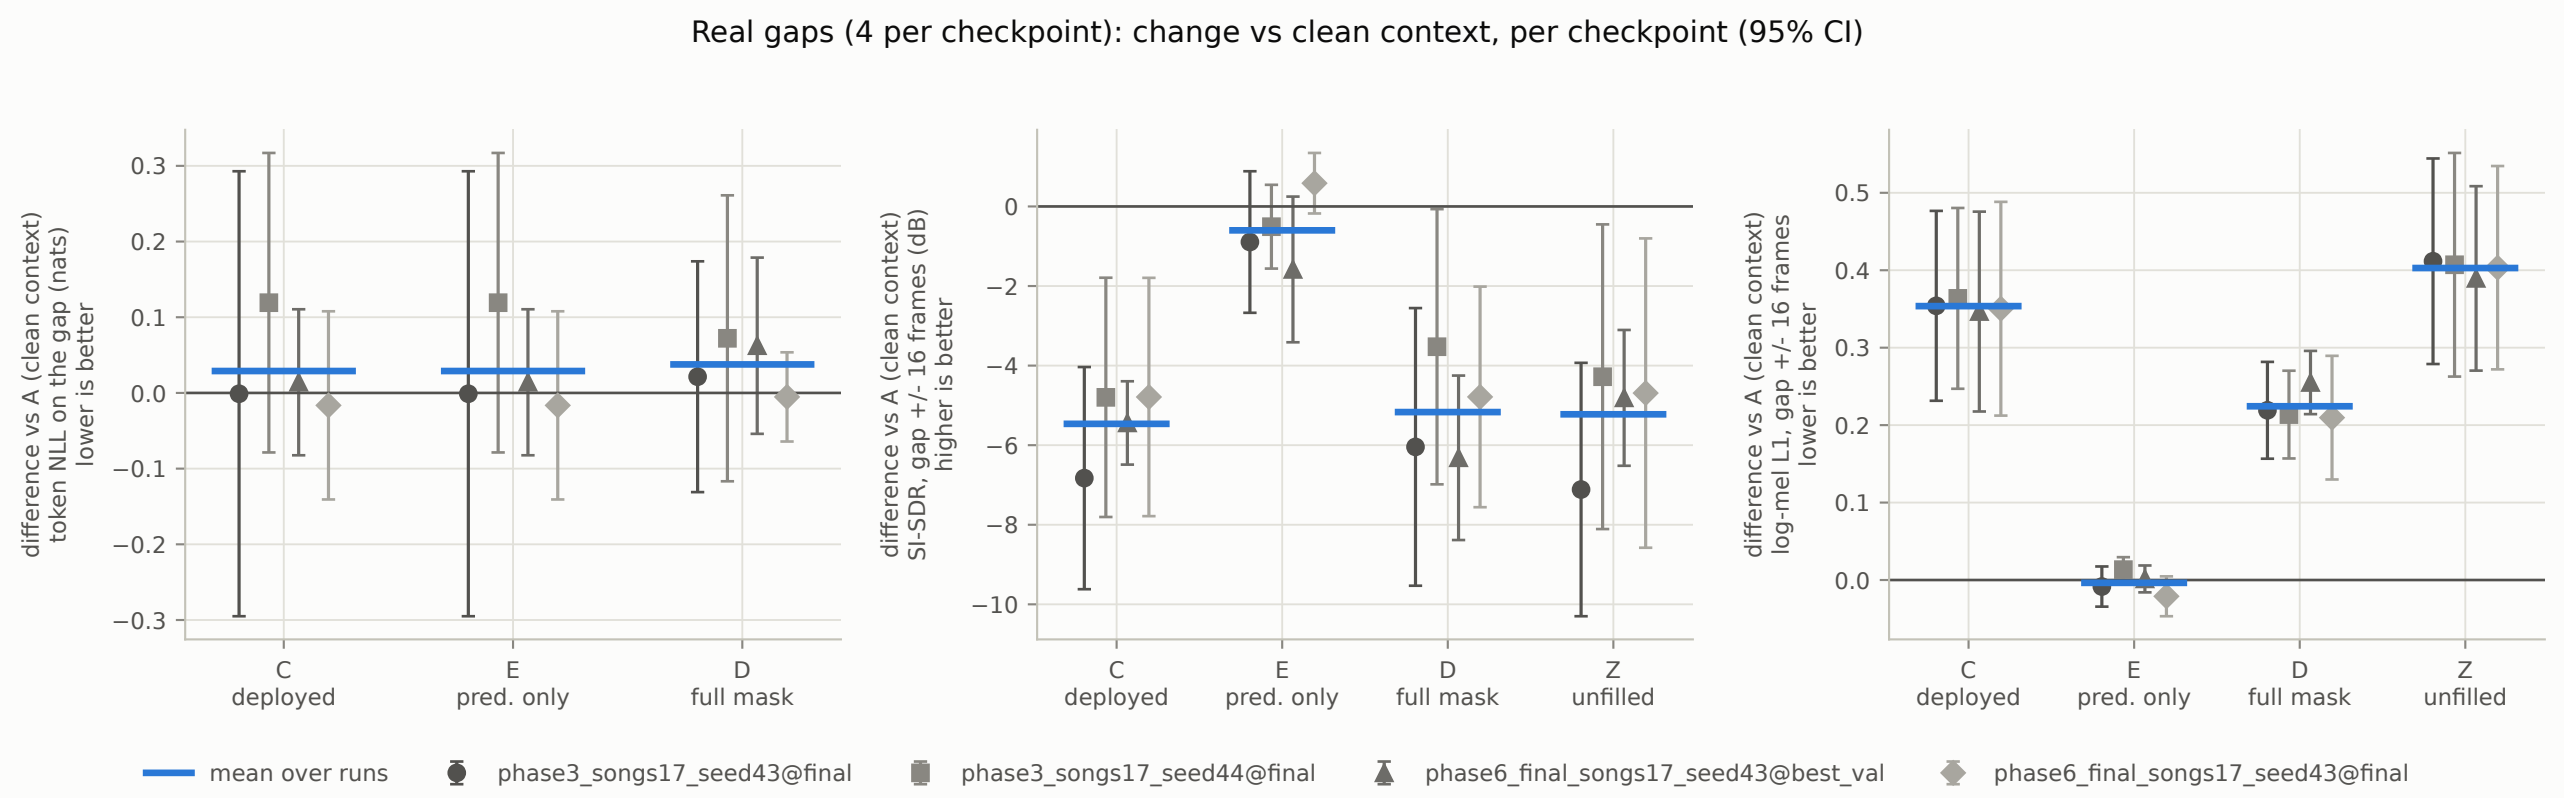
Real gaps (4 per checkpoint): change vs clean context, per checkpoint (95% CI)
difference vs A (clean context)
token NLL on the gap (nats)
lower is better
0.3
0.2
0.1
0.0
−0.1
−0.2
−0.3
C
deployed
E
pred. only
D
full mask
difference vs A (clean context)
SI-SDR, gap +/- 16 frames (dB)
higher is better
0
−2
−4
−6
−8
−10
C
deployed
E
pred. only
D
full mask
Z
unfilled
difference vs A (clean context)
log-mel L1, gap +/- 16 frames
lower is better
0.5
0.4
0.3
0.2
0.1
0.0
C
deployed
E
pred. only
D
full mask
Z
unfilled
mean over runs
phase3_songs17_seed43@final
phase3_songs17_seed44@final
phase6_final_songs17_seed43@best_val
phase6_final_songs17_seed43@final
A clean context (training-like) B frame-aligned zeroing C deployed today (unaligned zeroing, rounded mask)
E C's predictions written into clean context D mask covers every silent frame Z gap not inpainted